

ENGLISH

Français

Table of content

What is Named Entity Recognition (NER) ?

How Does Named Entity Recognition Work ?

Methods and Approaches for NER

Types of Named Entities

Key Challenges in Named Entity Recognition

Tools and Libraries for Named Entity Recognition

Practical Applications and Use Cases

Building a Named Entity Recognition Pipeline

Best Practices for Implementing NER

FAQ

Turning Text into Actionable Intelligence with NER

[Home](#) » [Blog](#) » [The Complete Guide to Named Entity Recognition \(NER\) : Methods, Tools, and Use Cases](#)

The Complete Guide

to Named Entity Recognition



The Complete Guide to Named Entity Recognition (NER): Methods, Tools, and Use Cases

April 16, 2025

Reading time: 10 min

Written by cnibart

Named Entity Recognition (NER) is a fundamental technique in Natural Language Processing (NLP) that involves identifying and classifying key elements, or "entities," within text into predefined categories such as names of persons, organizations, locations, dates, and more.

In this comprehensive guide, we'll delve into the intricacies of NER, exploring its underlying methodologies, the tools available for implementation, and the diverse applications it supports across various industries.

What is Named Entity Recognition (NER) ?

Definition and Purpose

Named Entity Recognition (NER) is a core NLP technique used to automatically detect and classify entities within unstructured text. These entities can represent people, organizations, locations, dates, quantities, or any other predefined category relevant to the analysis. The goal of NER is to convert raw textual data into structured information by identifying meaningful elements within the text. This allows downstream systems to better understand, search, and analyze language data.

AUTOMATICALLY FIND NAMES OF PEOPLE, PLACES, PRODUCTS, AND ORGANIZATIONS IN TEXT

How NER fits into Natural Language Processing (NLP) ?

In the NLP pipeline, NER typically comes after tasks like tokenization and part-of-speech (POS) tagging. It enriches the text by attaching semantic labels to words or phrases identified as entities. NER outputs can then be used by parsing systems, knowledge graphs, or information retrieval engines to enhance understanding and reasoning across documents.

Real-world Examples of Named Entities

- **Healthcare:** "Pfizer" (Organization), "COVID-19" (Medical term), "2020" (Date)
- **Legal:** "European Union" (Organization), "General Data Protection Regulation" (Law), "Paris" (Location)
- **Recruitment:** "Google" (Organization), "Data Scientist" (Job Title), "John Smith" (Person)

How Does Named Entity Recognition Work ?

The Process: From Tokenization to Classification

NER systems follow a structured pipeline to extract entities from raw text:

1. **Tokenization** – Splits the input text into individual words or tokens.
2. **POS Tagging** – Assigns a grammatical role (noun, verb, etc.) to each token.
3. **Entity Detection** – Identifies candidate tokens or spans likely to be entities.
4. **Entity Classification** – Labels each detected entity with a specific type (e.g., Person, Location, Organization).

This step-by-step flow transforms plain text into semantically enriched data that downstream applications can use for further analysis or decision-making.

Example of Annotated Text with NER Tags

Apple was founded by Steve Jobs in California in 1976.

Entity Recognition Output:

Text	Entity type
Apple	Organization
Steve Jobs	Person
California	Location
1976	Date

Using a trained model or a rule-augmented approach, our platform automatically tags and classifies entities, making unstructured text instantly searchable and ready for downstream business applications.

 Need domain-specific entities? Our interface lets you define and train custom entity types specific to your business needs — no coding required.

Methods and Approaches for NER

Rule-based Techniques

Rule-based systems rely on predefined patterns, such as regular expressions or curated dictionaries, to extract entities from text. For example, a rule like `r"\b[A-Z][a-z]+\b"` might capture person names like "John Smith". These methods are simple but brittle when handling ambiguity or unseen terms.

Machine Learning Models

Statistical models such as **Conditional Random Fields (CRF)** and **Support Vector Machines (SVM)** treat NER as a sequence labeling problem. Trained on annotated datasets, these models learn contextual patterns to predict entity boundaries and types, offering more adaptability than rule-based systems.

Deep Learning Approaches

Modern NER systems use neural networks like **BiLSTM** (Bidirectional Long Short-Term Memory) and **Transformers** to capture complex language features. These models can process long sequences, making them effective for identifying entities in unstructured, context-rich text.

Transfer Learning and BERT

Transfer learning leverages large pre-trained models like **BERT** (Bidirectional Encoder Representations from Transformers) fine-tuned on specific NER tasks. BERT-based NER systems achieve state-of-the-art accuracy by understanding nuanced language patterns without extensive task-specific training.

Hybrid Systems

Approach	Advantages	Limitations
----------	------------	-------------

ENGLISH Français		
BERT + Domain Rules	High precision in specific domains	Setup complexity, compute-intensive

Types of Named Entities

Generic Entity Categories (Person, Organization, etc.)

Standard NER systems typically identify the following core entity types:

- **Person** (e.g., "Marie Curie")
- **Organization** (e.g., "UNESCO")
- **Location** (e.g., "Tokyo")
- **Date** (e.g., "July 2021")
- **Product** (e.g., "iPhone")
- **Event** (e.g., "World Cup")

These categories form the baseline for many general-purpose information extraction tasks.

Domain-specific Entities

Domain	Example Entity Types
Healthcare	Drug name, Diagnosis, Procedure
Finance	Ticker, Currency, Financial instrument
Legal	Law reference, Jurisdiction, Contract clause
HR/Recruitment	Skill, Degree, Job title
Manufacturing	Part ID, Material, Machine type

Customizing entity types to suit the specific language and structure of a domain dramatically improves extraction quality and relevance.

Use in RAG chatbot

Kairntech-powered RAG chatbots leverage NER to enhance questions with structured context. For instance, when a user submits a query, the chatbot identifies and extracts key entities—such as product codes, project names, or client references—allowing the system to route the question to the most suitable agent for precise and efficient handling.

Key Challenges in Named Entity Recognition

Ambiguity and Context Dependency

Multilingual Issues

NER models trained in English don't generalize well to other languages. At Kairntech, we address this by supporting multilingual pipelines (e.g., English, French, German, Spanish, Dutch, Italian) and offering custom training for less-resourced languages through transfer learning.

Annotated Data Scarcity

High-quality training data is crucial but often lacking, especially in niche domains. Open datasets like WikiANN or CoNLL-2003 help, but domain-specific corpora still require manual annotation – a time-consuming process.

Domain Adaptation Difficulties

A model trained on news articles may fail on legal or technical documents. For example, "GAFA" in a tech context refers to organizations, but might go unrecognized in a general-purpose model. Adapting NER to specialized corpora requires custom training and iterative feedback loops – something our platform facilitates natively.



Tools and Libraries for Named Entity Recognition

spaCy Named Entity Recognition

Introduction to spaCy

spaCy is a fast, open-source NLP library in Python. It includes pre-trained NER models for several languages and supports deep learning integration out of the box.

How to Use spaCy for NER (Code Example)

```
import spacy

nlp = spacy.load("en_core_web_sm")

doc = nlp("Google acquired DeepMind in 2014.")
```

Google ORG

DeepMind ORG

2014 DATE

spaCy's pipeline automatically detects entities and assigns types such as organization and date using trained models.

Customizing spaCy for Your Domain

spaCy allows users to train or extend models with custom entity types using the EntityRuler or manual annotations. While powerful, this process still requires technical expertise and annotated data.

Other Popular Libraries

NLTK / Flair / Stanford NER

Library	Language	Usage Focus	Strengths
NLTK	Python	Educational / Prototyping	Lightweight, easy to start
Flair	Python	Deep Learning NER	Stacked embeddings, multilingual
Stanford NER	Java	Statistical NER	Reliable, mature models

Cloud-based APIs

Google / Amazon / IBM

- Google Cloud Natural Language
- Amazon Comprehend
- IBM Watson NLU

These services offer ready-to-use NER with scalable infrastructure but limited customization options.

Enterprise Solutions

How We at Kairntech Integrate NER ?

Our platform combines low-code interfaces with customizable NER models – supporting both standard and domain-specific entity types. Users can label data, train models, assess quality and deploy them securely, all within a no-install, enterprise-grade environment.

Practical Applications and Use Cases

Domain	How NER is Used
--------	-----------------

ENGLISH

Français

Biomedical Research

Identifies chemical compounds, and treatment entities in medical literature.

Legal Document Analysis

Detects contract clauses, legal terms, organization names, and jurisdiction references in case law.

Search Engines & Knowledge Graphs

Converts unstructured content into structured data to improve relevance and semantic linking.

Customer Service & Social Media

Tags product names, issues, locations, or sentiments in customer feedback for better response and analysis.

From healthcare to HR, NER supports scalable text analysis by converting language into structured, actionable information. At Kairntech, we help organizations leverage this power across domains with customizable assistants tailored to their data.

Building a Named Entity Recognition Pipeline

Data Collection and Annotation

The pipeline starts with gathering representative documents and annotating entities relevant to your use case — whether generic (like organization or person) or domain-specific (like part numbers or regulations).

Training and Evaluation

Annotated data is used to train a model, often via transfer learning. Evaluation follows using metrics such as precision, recall, and F1-score to ensure entity recognition quality aligns with business needs.

Deployment and Post-processing


After training, the model is integrated into an application or workflow. Post-processing steps—such as entity linking, normalization, or filtering—can be incorporated into the pipeline to ensure the outputs are business-ready and suitable for downstream use.

Continuous Feedback Loops

User corrections and new examples are reinjected into the system to retrain the model and improve accuracy over time — a key to maintaining performance in evolving environments.

Running Secure, On-premise NER with Kairntech

At Kairntech, our pipeline supports end-to-end NER — from raw document ingestion to entity extraction — entirely on-premise. This ensures data privacy while allowing teams to adapt models continuously, without writing code.



Best Practices for Implementing NER

Choosing the Right Model

1. Select a model architecture that fits your data scale and complexity – simple CRF models for small tasks, transformer-based models for high accuracy.

Handling Domain-specific Vocabulary

2. Use domain data to fine-tune models or enrich rule sets, ensuring accurate recognition of custom entity types not present in generic corpora.

Privacy and Security Considerations

3. Favor on-premise or private cloud deployments for sensitive information, especially in regulated industries like healthcare or law.

Empowering Teams with Low-code Tools

4. Enable subject matter experts to review, annotate, and improve models without coding – accelerating feedback cycles and improving outcomes.
5. Track performance over time (e.g., F1-score), version your models, and validate regularly to maintain consistency across use cases.

transformer-based models for high accuracy.

Handling Domain-specific Vocabulary

Use domain data to fine-tune models or enrich rule sets, ensuring accurate recognition of custom entity types not present in generic corpora.

Privacy and Security Considerations

Favor on-premise or private cloud deployments for sensitive information, especially in regulated industries like healthcare or law.

Empowering Teams with Low-code Tools

Once your hardware and software are in place, the next steps are installing, configuring and running the LLM.

Track performance over time (e.g., F1-score), version your models, and validate regularly to maintain consistency across use cases.

FAQ

How does an NER work ?

What is the difference between NER and NLP ?

What is the difference between BERT and NER ?

BERT is a large language model that can be fine-tuned for NER tasks. NER, on the other hand, is a goal — extracting entities — which BERT can help achieve when integrated into a recognition pipeline.

Is NER language-specific ?

Can NER detect emotions or sentiments ?

Turning Text into Actionable Intelligence with NER

The Future of Named Entity Recognition

ENGLISH

Français

How we at Kairntech Help Enterprises Build NER-Powered Assistants ?

We offer a secure, low-code platform to design, train, and deploy tailored NER solutions. Want to see it in action? Request a demo.

Share this

Can Kairntech's AI create business impact from your documents?

Book a call to find out →

← AI Document Analysis: Smarter Workflows with Intelligent Systems

Top Open-Source LLM Models in 2026: Features, Use Cases & Selection →

Related posts

ENGLISH

Français

How to Choose

The Best Text Annotation Tool for Your NLP Project ?



How to Choose the Best Text Annotation Tool for Your NLP Project ?

A text annotation tool is a software application that enables users to label specific elements...
August 5, 2026

FRANCE

29 Chemin du Vieux Chêne
38240 Meylan

GERMANY

Schützenhausstr. 28
69151 Neckargemünd

info@kairntech.com

ENGLISH

Français

Document class

Agentic RAG Chatbots

Semantic Enrichment

Pricing

Kairntech software

Kairntech consulting

Ressources

Blog

Manifesto

Documentation

REST API

Projects/Partnerships

Uses cases

Content discovery

Agentic RAG chatbot

Contact

Contact us

TERMS OF USE

GTC (SAAS)

GTC (LICENSE)

COOKIE POLICY (EU)

Kairntech © 2026 | Hosting: Pharéo